

Driver Drowsiness Detection using CNN



According to a report, around 40% of road accidents that happen on highways are caused by Drowsy Driving. This project aims at detecting whether a driver is feeling drowsy or is active while driving based on whether both the eyes of the driver are closed representing drowsiness or both the eyes are open using a CNN architecture.

Importing necessary libraries..

In [1]:

```
import numpy as np
import cv2
from tensorflow.keras.layers import Input, Conv2D, BatchNormalization, Dropout, Flatten, Dense, MaxPool2D
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.initializers import glorot_uniform
from tensorflow.keras.optimizers import Adam, SGD
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
from tensorflow.keras.regularizers import l2
import tensorflow as tf
import pandas as pd
from imutils.face_utils import FaceAligner
from imutils.face_utils import rect_to_bb
import imutils
import dlib
```

Loading the dataset and the test images

You can upload your own test images also.

In [6]:

```
!wget https://cainvas-static.s3.amazonaws.com/media/user_data/Yuvnish17/driver_drowsiness_detection_modified.zip
!unzip -qo driver_drowsiness_detection_modified.zip

--2021-07-03 15:56:18--  https://cainvas-static.s3.amazonaws.com/media/user_data/Yuvnish17/driver_drowsiness_detection_modified.zip
Resolving cainvas-static.s3.amazonaws.com (cainvas-static.s3.amazonaws.com)... 52.219.62.88
Connecting to cainvas-static.s3.amazonaws.com (cainvas-static.s3.amazonaws.com)|52.219.62.88|:443...
connected.
HTTP request sent, awaiting response... 200 OK
Length: 127436131 (122M) [application/x-zip-compressed]
Saving to: 'driver_drowsiness_detection_modified.zip'

driver_drowsiness_d 100%[=====>] 121.53M  98.2MB/s   in 1.2s

2021-07-03 15:56:19 (98.2 MB/s) - 'driver_drowsiness_detection_modified.zip' saved [127436131/127436131]
```

Loading the images and labels and Preprocessing the Dataset

Input Shape of Images for the CNN model - (32, 32, 3)

Dataset os split into Train Set and Test Set with test set containing 20% of the total number of images

Labels -> 0 - Closed Eye

1 - Open Eye

In [7]:

```
data = np.load('driver_drowsiness_detection/dataset_compressed.npz', allow_pickle=True)
X = data['arr_0']
Y = data['arr_1']

X = list(X)
Y = list(Y)
print(len(X))
print(len(Y))
```

1452

1452

In [8]:

```
for i in range(len(X)):
    img = X[i]
    img = cv2.resize(img, (32, 32))
    X[i] = img

print(len(X))
print(X[0].shape)
```

1452

(32, 32, 3)

In [9]:

```
label_encoder = LabelEncoder()
Y = label_encoder.fit_transform(Y)
print(Y.shape)
print(Y[0])
print(set(Y))
```

(1452,)

0

{0, 1}

In [10]:

```
X = np.array(X)
Y = np.array(Y)
print(X.shape)
print(Y.shape)
```

(1452, 32, 32, 3)

(1452,)

In [11]:

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2)
```

In [12]:

```
print(X_train.shape)
print(Y_train.shape)
print(X_test.shape)
print(Y_test.shape)

Y_train = to_categorical(Y_train)
Y_test = to_categorical(Y_test)

print(X_train.shape)
print(Y_train.shape)
print(X_test.shape)
print(Y_test.shape)
```

(1161, 32, 32, 3)

(1161,)

(291, 32, 32, 3)

(291,)

(1161, 32, 32, 3)

(1161, 2)

(291, 32, 32, 3)

(291, 2)

Visualizing images of Closed Eye and Open Eye from the Dataset

In [13]:

```
figure1 = plt.figure(figsize=(5, 5))
idx_closed = np.where(Y==0)
img_closed = X[idx_closed[0][0]]
plt.imshow(img_closed)
plt.title('Image of Closed Eye representing Driver is sleeping')
plt.axis('off')
plt.show()

figure2 = plt.figure(figsize=(5, 5))
idx_open = np.where(Y==1)
img_open = X[idx_open[0][0]]
plt.imshow(img_open)
plt.title('Image of Open Eye representing Driver is not sleeping')
plt.axis('off')
plt.show()
```

Image of Closed Eye representing Driver is sleeping



Image of Open Eye representing Driver is not sleeping




Visualizing the Data Distribution in Train and Test Set

In [11]:

```
unique_train, count = np.unique(Y_train, return_counts=True)
plt.figure(figsize=(20, 10))
sns.barplot(unique_train, count).set_title("Number of training images per category:")
plt.show()
```


/opt/tljh/user/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning



In [12]:

```
unique_test, count_test = np.unique(Y_test, return_counts=True)
plt.figure(figsize=(20, 10))
sns.barplot(unique_test, count_test).set_title("Number of training images per category:")
plt.show()
```

/opt/tljh/user/lib/python3.7/site-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning



As can be seen from the bar plots that the dataset is already balanced so no need to perform dataset balancing

Defining the Model Architecture

In [13]:

```
def driver_drowsiness_detection_model(input_shape=(32, 32, 3)):
    model = Sequential()
    model.add(Input(shape=input_shape))
    model.add(Conv2D(32, (3, 3), padding='same', strides=(1, 1), name='conv1', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(Conv2D(32, (3, 3), padding='same', strides=(1, 1), name='conv2', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(Dropout(0.2))
    model.add(MaxPool2D((2, 2), strides=(2, 2)))

    model.add(Conv2D(64, (3, 3), padding='same', strides=(1, 1), name='conv3', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(MaxPool2D((2, 2), strides=(2, 2)))
    model.add(Conv2D(64, (3, 3), padding='same', strides=(1, 1), name='conv4', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(Dropout(0.3))
    model.add(MaxPool2D((2, 2), strides=(2, 2)))

    model.add(Conv2D(64, (3, 3), padding='same', strides=(1, 1), name='conv5', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(Conv2D(64, (3, 3), padding='same', strides=(1, 1), name='conv6', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(Conv2D(64, (3, 3), padding='same', strides=(1, 1), name='conv7', activation='relu',
                     kernel_initializer=glorot_uniform(seed=0)))
    model.add(BatchNormalization())
    model.add(Dropout(0.4))
    model.add(MaxPool2D((2, 2), strides=(2, 2)))

    model.add(Flatten())
    model.add(Dense(128, activation='relu', kernel_initializer=glorot_uniform(seed=0), name='fc1'))
    model.add(Dropout(0.5))
    model.add(Dense(128, activation='relu', kernel_initializer=glorot_uniform(seed=0), name='fc2'))
    model.add(Dropout(0.5))
    model.add(Dense(2, activation='softmax', kernel_initializer=glorot_uniform(seed=0), name='fc3'))

    optimizer = Adam(0.0001)
    model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
    return model
```

In [14]:

```
model= driver_drowsiness_detection_model(input_shape=(32, 32, 3))
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (Batch Normalization)	(None, 32, 32, 32)	128
conv2 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 32)	128
dropout (Dropout)	(None, 32, 32, 32)	0
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv3 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_2 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv4 (Conv2D)	(None, 8, 8, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 8, 8, 64)	256
dropout_1 (Dropout)	(None, 8, 8, 64)	0
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
conv5 (Conv2D)	(None, 4, 4, 64)	36928
batch_normalization_4 (Batch Normalization)	(None, 4, 4, 64)	256
conv6 (Conv2D)	(None, 4, 4, 64)	36928
batch_normalization_5 (Batch Normalization)	(None, 4, 4, 64)	256
conv7 (Conv2D)	(None, 4, 4, 64)	36928
batch_normalization_6 (Batch Normalization)	(None, 4, 4, 64)	256
dropout_2 (Dropout)	(None, 4, 4, 64)	0
max_pooling2d_3 (MaxPooling2D)	(None, 2, 2, 64)	0
flatten (Flatten)	(None, 256)	0
fc1 (Dense)	(None, 128)	32896
dropout_3 (Dropout)	(None, 128)	0
fc2 (Dense)	(None, 128)	16512
dropout_4 (Dropout)	(None, 128)	0
fc3 (Dense)	(None, 2)	258
=====		
Total params: 227,554		
Trainable params: 226,786		
Non-trainable params: 768		

Training the Model

Learning Rate - 0.0001
Optimizer - Adam
Number of Epochs - 200
Batch Size = 128

In [15]:

```
aug = ImageDataGenerator(rotation_range=20, zoom_range=0.2, horizontal_flip=True)
hist = model.fit(aug.flow(X_train, Y_train, batch_size=128), batch_size=128, epochs=200, validation_data=(X_test, Y_test))
```

```
Epoch 1/200
10/10 [=====] - 1s 62ms/step - loss: 1.3128 - accuracy: 0.5357 - val_loss: 0.7441 -
val_accuracy: 0.4192
Epoch 2/200
10/10 [=====] - 0s 35ms/step - loss: 1.0820 - accuracy: 0.5719 - val_loss: 0.7566 -
val_accuracy: 0.4296
Epoch 3/200
10/10 [=====] - 0s 35ms/step - loss: 1.0270 - accuracy: 0.5461 - val_loss: 0.7175 -
val_accuracy: 0.4880
Epoch 4/200
10/10 [=====] - 0s 35ms/step - loss: 0.9197 - accuracy: 0.5642 - val_loss: 0.6779 -
val_accuracy: 0.5464
Epoch 5/200
10/10 [=====] - 0s 39ms/step - loss: 0.8398 - accuracy: 0.6038 - val_loss: 0.6467 -
val_accuracy: 0.6117
Epoch 6/200
10/10 [=====] - 0s 39ms/step - loss: 0.8226 - accuracy: 0.5848 - val_loss: 0.6234 -
val_accuracy: 0.6907
Epoch 7/200
10/10 [=====] - 0s 35ms/step - loss: 0.8189 - accuracy: 0.6012 - val_loss: 0.6081 -
val_accuracy: 0.7320
Epoch 8/200
10/10 [=====] - 0s 35ms/step - loss: 0.7352 - accuracy: 0.6202 - val_loss: 0.5990 -
val_accuracy: 0.7904
Epoch 9/200
10/10 [=====] - 0s 36ms/step - loss: 0.6566 - accuracy: 0.6520 - val_loss: 0.5902 -
val_accuracy: 0.7938
Epoch 10/200
10/10 [=====] - 0s 36ms/step - loss: 0.7122 - accuracy: 0.6417 - val_loss: 0.5802 -
val_accuracy: 0.7835
Epoch 11/200
10/10 [=====] - 0s 36ms/step - loss: 0.6710 - accuracy: 0.6503 - val_loss: 0.5703 -
val_accuracy: 0.7869
Epoch 12/200
10/10 [=====] - 0s 36ms/step - loss: 0.7311 - accuracy: 0.6348 - val_loss: 0.5607 -
val_accuracy: 0.8007
Epoch 13/200
10/10 [=====] - 0s 35ms/step - loss: 0.6478 - accuracy: 0.6555 - val_loss: 0.5492 -
val_accuracy: 0.8007
Epoch 14/200
10/10 [=====] - 0s 39ms/step - loss: 0.6640 - accuracy: 0.6770 - val_loss: 0.5391 -
val_accuracy: 0.8041
Epoch 15/200
10/10 [=====] - 0s 37ms/step - loss: 0.6336 - accuracy: 0.6882 - val_loss: 0.5273 -
val_accuracy: 0.8076
Epoch 16/200
10/10 [=====] - 0s 36ms/step - loss: 0.6174 - accuracy: 0.6848 - val_loss: 0.5130 -
val_accuracy: 0.8144
Epoch 17/200
10/10 [=====] - 0s 36ms/step - loss: 0.5974 - accuracy: 0.6951 - val_loss: 0.4972 -
val_accuracy: 0.8179
Epoch 18/200
10/10 [=====] - 0s 36ms/step - loss: 0.6010 - accuracy: 0.7054 - val_loss: 0.4797 -
val_accuracy: 0.8351
Epoch 19/200
10/10 [=====] - 0s 36ms/step - loss: 0.5604 - accuracy: 0.7140 - val_loss: 0.4619 -
val_accuracy: 0.8591
Epoch 20/200
10/10 [=====] - 0s 36ms/step - loss: 0.6022 - accuracy: 0.7115 - val_loss: 0.4411 -
val_accuracy: 0.8660
Epoch 21/200
10/10 [=====] - 0s 36ms/step - loss: 0.5179 - accuracy: 0.7407 - val_loss: 0.4199 -
val_accuracy: 0.8797
Epoch 22/200
10/10 [=====] - 0s 37ms/step - loss: 0.5241 - accuracy: 0.7571 - val_loss: 0.3991 -
val_accuracy: 0.8832
Epoch 23/200
10/10 [=====] - 0s 36ms/step - loss: 0.5258 - accuracy: 0.7485 - val_loss: 0.3811 -
val_accuracy: 0.8866
Epoch 24/200
10/10 [=====] - 0s 36ms/step - loss: 0.5202 - accuracy: 0.7657 - val_loss: 0.3606 -
val_accuracy: 0.8969
Epoch 25/200
10/10 [=====] - 0s 37ms/step - loss: 0.5184 - accuracy: 0.7502 - val_loss: 0.3444 -
val_accuracy: 0.9003
Epoch 26/200
10/10 [=====] - 0s 36ms/step - loss: 0.4973 - accuracy: 0.7580 - val_loss: 0.3311 -
val_accuracy: 0.9210
Epoch 27/200
10/10 [=====] - 0s 36ms/step - loss: 0.5014 - accuracy: 0.7692 - val_loss: 0.3167 -
val_accuracy: 0.9244
Epoch 28/200
10/10 [=====] - 0s 36ms/step - loss: 0.4590 - accuracy: 0.7778 - val_loss: 0.3032 -
val_accuracy: 0.9278
Epoch 29/200
10/10 [=====] - 0s 37ms/step - loss: 0.4679 - accuracy: 0.7786 - val_loss: 0.2884 -
val_accuracy: 0.9313
Epoch 30/200
```



```
10/10 [=====] - 0s 36ms/step - loss: 0.4664 - accuracy: 0.7933 - val_loss: 0.2724 -  
val_accuracy: 0.9381  
Epoch 31/200  
10/10 [=====] - 0s 36ms/step - loss: 0.4302 - accuracy: 0.7924 - val_loss: 0.2581 -  
val_accuracy: 0.9416  
Epoch 32/200  
10/10 [=====] - 0s 37ms/step - loss: 0.4441 - accuracy: 0.8174 - val_loss: 0.2453 -  
val_accuracy: 0.9381  
Epoch 33/200  
10/10 [=====] - 0s 36ms/step - loss: 0.3975 - accuracy: 0.8329 - val_loss: 0.2322 -  
val_accuracy: 0.9381  
Epoch 34/200  
10/10 [=====] - 0s 37ms/step - loss: 0.4253 - accuracy: 0.8243 - val_loss: 0.2171 -  
val_accuracy: 0.9381  
Epoch 35/200  
10/10 [=====] - 0s 36ms/step - loss: 0.4118 - accuracy: 0.8208 - val_loss: 0.2050 -  
val_accuracy: 0.9381  
Epoch 36/200  
10/10 [=====] - 0s 37ms/step - loss: 0.3658 - accuracy: 0.8303 - val_loss: 0.1964 -  
val_accuracy: 0.9347  
Epoch 37/200  
10/10 [=====] - 0s 37ms/step - loss: 0.3687 - accuracy: 0.8303 - val_loss: 0.1888 -  
val_accuracy: 0.9313  
Epoch 38/200  
10/10 [=====] - 0s 36ms/step - loss: 0.3588 - accuracy: 0.8519 - val_loss: 0.1797 -  
val_accuracy: 0.9347  
Epoch 39/200  
10/10 [=====] - 0s 36ms/step - loss: 0.3430 - accuracy: 0.8587 - val_loss: 0.1735 -  
val_accuracy: 0.9381  
Epoch 40/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2976 - accuracy: 0.8682 - val_loss: 0.1685 -  
val_accuracy: 0.9313  
Epoch 41/200  
10/10 [=====] - 0s 36ms/step - loss: 0.3316 - accuracy: 0.8579 - val_loss: 0.1632 -  
val_accuracy: 0.9347  
Epoch 42/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2968 - accuracy: 0.8811 - val_loss: 0.1543 -  
val_accuracy: 0.9381  
Epoch 43/200  
10/10 [=====] - 0s 36ms/step - loss: 0.3112 - accuracy: 0.8760 - val_loss: 0.1443 -  
val_accuracy: 0.9485  
Epoch 44/200  
10/10 [=====] - 0s 35ms/step - loss: 0.2778 - accuracy: 0.8880 - val_loss: 0.1405 -  
val_accuracy: 0.9485  
Epoch 45/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2871 - accuracy: 0.8820 - val_loss: 0.1415 -  
val_accuracy: 0.9381  
Epoch 46/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2912 - accuracy: 0.8889 - val_loss: 0.1260 -  
val_accuracy: 0.9656  
Epoch 47/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2999 - accuracy: 0.8725 - val_loss: 0.1177 -  
val_accuracy: 0.9656  
Epoch 48/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2901 - accuracy: 0.8811 - val_loss: 0.1133 -  
val_accuracy: 0.9622  
Epoch 49/200  
10/10 [=====] - 0s 37ms/step - loss: 0.2573 - accuracy: 0.9061 - val_loss: 0.1134 -  
val_accuracy: 0.9622  
Epoch 50/200  
10/10 [=====] - 0s 35ms/step - loss: 0.2705 - accuracy: 0.8949 - val_loss: 0.1095 -  
val_accuracy: 0.9622  
Epoch 51/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2763 - accuracy: 0.9001 - val_loss: 0.1096 -  
val_accuracy: 0.9656  
Epoch 52/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2430 - accuracy: 0.9104 - val_loss: 0.1019 -  
val_accuracy: 0.9691  
Epoch 53/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2337 - accuracy: 0.9078 - val_loss: 0.1000 -  
val_accuracy: 0.9656  
Epoch 54/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2363 - accuracy: 0.9027 - val_loss: 0.0948 -  
val_accuracy: 0.9691  
Epoch 55/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2478 - accuracy: 0.9165 - val_loss: 0.0898 -  
val_accuracy: 0.9725  
Epoch 56/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2253 - accuracy: 0.9070 - val_loss: 0.0872 -  
val_accuracy: 0.9725  
Epoch 57/200  
10/10 [=====] - 0s 37ms/step - loss: 0.2045 - accuracy: 0.9139 - val_loss: 0.0845 -  
val_accuracy: 0.9725  
Epoch 58/200  
10/10 [=====] - 0s 37ms/step - loss: 0.2179 - accuracy: 0.9147 - val_loss: 0.0821 -  
val_accuracy: 0.9691  
Epoch 59/200  
10/10 [=====] - 0s 36ms/step - loss: 0.2063 - accuracy: 0.9268 - val_loss: 0.0776 -  
val_accuracy: 0.9759
```

```
Epoch 60/200
10/10 [=====] - 0s 36ms/step - loss: 0.1910 - accuracy: 0.9251 - val_loss: 0.0774 -
val_accuracy: 0.9759
Epoch 61/200
10/10 [=====] - 0s 36ms/step - loss: 0.2069 - accuracy: 0.9087 - val_loss: 0.0789 -
val_accuracy: 0.9725
Epoch 62/200
10/10 [=====] - 0s 36ms/step - loss: 0.2207 - accuracy: 0.9182 - val_loss: 0.0767 -
val_accuracy: 0.9725
Epoch 63/200
10/10 [=====] - 0s 36ms/step - loss: 0.1807 - accuracy: 0.9242 - val_loss: 0.0803 -
val_accuracy: 0.9725
Epoch 64/200
10/10 [=====] - 0s 36ms/step - loss: 0.1843 - accuracy: 0.9320 - val_loss: 0.0843 -
val_accuracy: 0.9691
Epoch 65/200
10/10 [=====] - 0s 39ms/step - loss: 0.1745 - accuracy: 0.9388 - val_loss: 0.0890 -
val_accuracy: 0.9656
Epoch 66/200
10/10 [=====] - 0s 36ms/step - loss: 0.1799 - accuracy: 0.9320 - val_loss: 0.0809 -
val_accuracy: 0.9656
Epoch 67/200
10/10 [=====] - 0s 36ms/step - loss: 0.1914 - accuracy: 0.9285 - val_loss: 0.0788 -
val_accuracy: 0.9691
Epoch 68/200
10/10 [=====] - 0s 37ms/step - loss: 0.1768 - accuracy: 0.9302 - val_loss: 0.0685 -
val_accuracy: 0.9656
Epoch 69/200
10/10 [=====] - 0s 39ms/step - loss: 0.1577 - accuracy: 0.9414 - val_loss: 0.0663 -
val_accuracy: 0.9759
Epoch 70/200
10/10 [=====] - 0s 36ms/step - loss: 0.1637 - accuracy: 0.9432 - val_loss: 0.0640 -
val_accuracy: 0.9759
Epoch 71/200
10/10 [=====] - 0s 37ms/step - loss: 0.1806 - accuracy: 0.9320 - val_loss: 0.0687 -
val_accuracy: 0.9759
Epoch 72/200
10/10 [=====] - 0s 39ms/step - loss: 0.1754 - accuracy: 0.9294 - val_loss: 0.0897 -
val_accuracy: 0.9691
Epoch 73/200
10/10 [=====] - 0s 37ms/step - loss: 0.1697 - accuracy: 0.9423 - val_loss: 0.0744 -
val_accuracy: 0.9759
Epoch 74/200
10/10 [=====] - 0s 39ms/step - loss: 0.1511 - accuracy: 0.9492 - val_loss: 0.0717 -
val_accuracy: 0.9691
Epoch 75/200
10/10 [=====] - 0s 36ms/step - loss: 0.1606 - accuracy: 0.9388 - val_loss: 0.0730 -
val_accuracy: 0.9691
Epoch 76/200
10/10 [=====] - 0s 36ms/step - loss: 0.1564 - accuracy: 0.9397 - val_loss: 0.0731 -
val_accuracy: 0.9691
Epoch 77/200
10/10 [=====] - 0s 36ms/step - loss: 0.1386 - accuracy: 0.9518 - val_loss: 0.0745 -
val_accuracy: 0.9622
Epoch 78/200
10/10 [=====] - 0s 36ms/step - loss: 0.1787 - accuracy: 0.9414 - val_loss: 0.0758 -
val_accuracy: 0.9725
Epoch 79/200
10/10 [=====] - 0s 35ms/step - loss: 0.1468 - accuracy: 0.9535 - val_loss: 0.0742 -
val_accuracy: 0.9725
Epoch 80/200
10/10 [=====] - 0s 39ms/step - loss: 0.1503 - accuracy: 0.9449 - val_loss: 0.0810 -
val_accuracy: 0.9725
Epoch 81/200
10/10 [=====] - 0s 36ms/step - loss: 0.1587 - accuracy: 0.9492 - val_loss: 0.0822 -
val_accuracy: 0.9725
Epoch 82/200
10/10 [=====] - 0s 35ms/step - loss: 0.1370 - accuracy: 0.9483 - val_loss: 0.0816 -
val_accuracy: 0.9691
Epoch 83/200
10/10 [=====] - 0s 37ms/step - loss: 0.1517 - accuracy: 0.9440 - val_loss: 0.0829 -
val_accuracy: 0.9725
Epoch 84/200
10/10 [=====] - 0s 37ms/step - loss: 0.1514 - accuracy: 0.9440 - val_loss: 0.0955 -
val_accuracy: 0.9588
Epoch 85/200
10/10 [=====] - 0s 35ms/step - loss: 0.1516 - accuracy: 0.9363 - val_loss: 0.0848 -
val_accuracy: 0.9691
Epoch 86/200
10/10 [=====] - 0s 36ms/step - loss: 0.1268 - accuracy: 0.9543 - val_loss: 0.0862 -
val_accuracy: 0.9691
Epoch 87/200
10/10 [=====] - 0s 35ms/step - loss: 0.1279 - accuracy: 0.9518 - val_loss: 0.0752 -
val_accuracy: 0.9656
Epoch 88/200
10/10 [=====] - 0s 36ms/step - loss: 0.1170 - accuracy: 0.9561 - val_loss: 0.0740 -
val_accuracy: 0.9691
Epoch 89/200
10/10 [=====] - 0s 40ms/step - loss: 0.1306 - accuracy: 0.9483 - val_loss: 0.0781 -
```

```
val_accuracy: 0.9656
Epoch 90/200
10/10 [=====] - 0s 36ms/step - loss: 0.1162 - accuracy: 0.9535 - val_loss: 0.0853 -
val_accuracy: 0.9691
Epoch 91/200
10/10 [=====] - 0s 35ms/step - loss: 0.1307 - accuracy: 0.9535 - val_loss: 0.0746 -
val_accuracy: 0.9691
Epoch 92/200
10/10 [=====] - 0s 36ms/step - loss: 0.1218 - accuracy: 0.9595 - val_loss: 0.0618 -
val_accuracy: 0.9759
Epoch 93/200
10/10 [=====] - 0s 36ms/step - loss: 0.1346 - accuracy: 0.9569 - val_loss: 0.0625 -
val_accuracy: 0.9759
Epoch 94/200
10/10 [=====] - 0s 39ms/step - loss: 0.1269 - accuracy: 0.9569 - val_loss: 0.0698 -
val_accuracy: 0.9656
Epoch 95/200
10/10 [=====] - 0s 36ms/step - loss: 0.1172 - accuracy: 0.9587 - val_loss: 0.0594 -
val_accuracy: 0.9759
Epoch 96/200
10/10 [=====] - 0s 39ms/step - loss: 0.1206 - accuracy: 0.9569 - val_loss: 0.0669 -
val_accuracy: 0.9691
Epoch 97/200
10/10 [=====] - 0s 37ms/step - loss: 0.1250 - accuracy: 0.9483 - val_loss: 0.0695 -
val_accuracy: 0.9691
Epoch 98/200
10/10 [=====] - 0s 36ms/step - loss: 0.1014 - accuracy: 0.9699 - val_loss: 0.0600 -
val_accuracy: 0.9759
Epoch 99/200
10/10 [=====] - 0s 36ms/step - loss: 0.1229 - accuracy: 0.9569 - val_loss: 0.0635 -
val_accuracy: 0.9759
Epoch 100/200
10/10 [=====] - 0s 35ms/step - loss: 0.0951 - accuracy: 0.9681 - val_loss: 0.0698 -
val_accuracy: 0.9725
Epoch 101/200
10/10 [=====] - 0s 37ms/step - loss: 0.0858 - accuracy: 0.9673 - val_loss: 0.0672 -
val_accuracy: 0.9725
Epoch 102/200
10/10 [=====] - 0s 36ms/step - loss: 0.0766 - accuracy: 0.9707 - val_loss: 0.0669 -
val_accuracy: 0.9725
Epoch 103/200
10/10 [=====] - 0s 36ms/step - loss: 0.0794 - accuracy: 0.9673 - val_loss: 0.0670 -
val_accuracy: 0.9691
Epoch 104/200
10/10 [=====] - 0s 37ms/step - loss: 0.1060 - accuracy: 0.9647 - val_loss: 0.0626 -
val_accuracy: 0.9691
Epoch 105/200
10/10 [=====] - 0s 40ms/step - loss: 0.1153 - accuracy: 0.9612 - val_loss: 0.0710 -
val_accuracy: 0.9794
Epoch 106/200
10/10 [=====] - 0s 37ms/step - loss: 0.1148 - accuracy: 0.9595 - val_loss: 0.0701 -
val_accuracy: 0.9794
Epoch 107/200
10/10 [=====] - 0s 36ms/step - loss: 0.1073 - accuracy: 0.9595 - val_loss: 0.0705 -
val_accuracy: 0.9794
Epoch 108/200
10/10 [=====] - 0s 36ms/step - loss: 0.0861 - accuracy: 0.9638 - val_loss: 0.0686 -
val_accuracy: 0.9794
Epoch 109/200
10/10 [=====] - 0s 36ms/step - loss: 0.0800 - accuracy: 0.9716 - val_loss: 0.0691 -
val_accuracy: 0.9828
Epoch 110/200
10/10 [=====] - 0s 35ms/step - loss: 0.0823 - accuracy: 0.9699 - val_loss: 0.0672 -
val_accuracy: 0.9794
Epoch 111/200
10/10 [=====] - 0s 39ms/step - loss: 0.0861 - accuracy: 0.9716 - val_loss: 0.0681 -
val_accuracy: 0.9794
Epoch 112/200
10/10 [=====] - 0s 36ms/step - loss: 0.0984 - accuracy: 0.9664 - val_loss: 0.0672 -
val_accuracy: 0.9759
Epoch 113/200
10/10 [=====] - 0s 36ms/step - loss: 0.0905 - accuracy: 0.9664 - val_loss: 0.0785 -
val_accuracy: 0.9725
Epoch 114/200
10/10 [=====] - 0s 36ms/step - loss: 0.1056 - accuracy: 0.9587 - val_loss: 0.0893 -
val_accuracy: 0.9691
Epoch 115/200
10/10 [=====] - 0s 37ms/step - loss: 0.0725 - accuracy: 0.9742 - val_loss: 0.0798 -
val_accuracy: 0.9725
Epoch 116/200
10/10 [=====] - 0s 36ms/step - loss: 0.1012 - accuracy: 0.9612 - val_loss: 0.0737 -
val_accuracy: 0.9759
Epoch 117/200
10/10 [=====] - 0s 36ms/step - loss: 0.0977 - accuracy: 0.9587 - val_loss: 0.0763 -
val_accuracy: 0.9794
Epoch 118/200
10/10 [=====] - 0s 36ms/step - loss: 0.0768 - accuracy: 0.9724 - val_loss: 0.0882 -
val_accuracy: 0.9725
Epoch 119/200
```

```
10/10 [=====] - 0s 36ms/step - loss: 0.1243 - accuracy: 0.9587 - val_loss: 0.0810 -  
val_accuracy: 0.9725  
Epoch 120/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0814 - accuracy: 0.9742 - val_loss: 0.0868 -  
val_accuracy: 0.9725  
Epoch 121/200  
10/10 [=====] - 0s 37ms/step - loss: 0.1041 - accuracy: 0.9578 - val_loss: 0.0934 -  
val_accuracy: 0.9794  
Epoch 122/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0879 - accuracy: 0.9664 - val_loss: 0.0962 -  
val_accuracy: 0.9656  
Epoch 123/200  
10/10 [=====] - 0s 37ms/step - loss: 0.0795 - accuracy: 0.9767 - val_loss: 0.0838 -  
val_accuracy: 0.9691  
Epoch 124/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0665 - accuracy: 0.9742 - val_loss: 0.0691 -  
val_accuracy: 0.9794  
Epoch 125/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0682 - accuracy: 0.9690 - val_loss: 0.0729 -  
val_accuracy: 0.9794  
Epoch 126/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0872 - accuracy: 0.9647 - val_loss: 0.0761 -  
val_accuracy: 0.9794  
Epoch 127/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0846 - accuracy: 0.9750 - val_loss: 0.0664 -  
val_accuracy: 0.9828  
Epoch 128/200  
10/10 [=====] - 0s 37ms/step - loss: 0.0676 - accuracy: 0.9802 - val_loss: 0.0704 -  
val_accuracy: 0.9828  
Epoch 129/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0813 - accuracy: 0.9690 - val_loss: 0.0807 -  
val_accuracy: 0.9828  
Epoch 130/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0811 - accuracy: 0.9733 - val_loss: 0.0876 -  
val_accuracy: 0.9691  
Epoch 131/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0766 - accuracy: 0.9733 - val_loss: 0.0604 -  
val_accuracy: 0.9828  
Epoch 132/200  
10/10 [=====] - 0s 39ms/step - loss: 0.0961 - accuracy: 0.9681 - val_loss: 0.0677 -  
val_accuracy: 0.9759  
Epoch 133/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0843 - accuracy: 0.9681 - val_loss: 0.0988 -  
val_accuracy: 0.9725  
Epoch 134/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0771 - accuracy: 0.9742 - val_loss: 0.1048 -  
val_accuracy: 0.9691  
Epoch 135/200  
10/10 [=====] - 0s 40ms/step - loss: 0.0692 - accuracy: 0.9776 - val_loss: 0.0993 -  
val_accuracy: 0.9656  
Epoch 136/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0880 - accuracy: 0.9733 - val_loss: 0.1205 -  
val_accuracy: 0.9588  
Epoch 137/200  
10/10 [=====] - 0s 37ms/step - loss: 0.0690 - accuracy: 0.9733 - val_loss: 0.1100 -  
val_accuracy: 0.9553  
Epoch 138/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0847 - accuracy: 0.9707 - val_loss: 0.1003 -  
val_accuracy: 0.9588  
Epoch 139/200  
10/10 [=====] - 0s 35ms/step - loss: 0.0715 - accuracy: 0.9750 - val_loss: 0.0811 -  
val_accuracy: 0.9622  
Epoch 140/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0700 - accuracy: 0.9742 - val_loss: 0.0743 -  
val_accuracy: 0.9622  
Epoch 141/200  
10/10 [=====] - 0s 37ms/step - loss: 0.0758 - accuracy: 0.9750 - val_loss: 0.0739 -  
val_accuracy: 0.9622  
Epoch 142/200  
10/10 [=====] - 0s 40ms/step - loss: 0.0972 - accuracy: 0.9655 - val_loss: 0.0750 -  
val_accuracy: 0.9622  
Epoch 143/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0758 - accuracy: 0.9742 - val_loss: 0.0708 -  
val_accuracy: 0.9656  
Epoch 144/200  
10/10 [=====] - 0s 43ms/step - loss: 0.0713 - accuracy: 0.9716 - val_loss: 0.0593 -  
val_accuracy: 0.9759  
Epoch 145/200  
10/10 [=====] - 0s 45ms/step - loss: 0.0555 - accuracy: 0.9802 - val_loss: 0.0629 -  
val_accuracy: 0.9794  
Epoch 146/200  
10/10 [=====] - 0s 36ms/step - loss: 0.0542 - accuracy: 0.9819 - val_loss: 0.0627 -  
val_accuracy: 0.9828  
Epoch 147/200  
10/10 [=====] - 0s 37ms/step - loss: 0.0531 - accuracy: 0.9871 - val_loss: 0.0663 -  
val_accuracy: 0.9794  
Epoch 148/200  
10/10 [=====] - 0s 37ms/step - loss: 0.0500 - accuracy: 0.9819 - val_loss: 0.0635 -  
val_accuracy: 0.9794
```

Epoch 149/200
10/10 [=====] - 0s 39ms/step - loss: 0.0457 - accuracy: 0.9845 - val_loss: 0.0704 - val_accuracy: 0.9794
Epoch 150/200
10/10 [=====] - 0s 36ms/step - loss: 0.0583 - accuracy: 0.9793 - val_loss: 0.0874 - val_accuracy: 0.9656
Epoch 151/200
10/10 [=====] - 0s 37ms/step - loss: 0.0732 - accuracy: 0.9793 - val_loss: 0.0891 - val_accuracy: 0.9656
Epoch 152/200
10/10 [=====] - 0s 36ms/step - loss: 0.0553 - accuracy: 0.9811 - val_loss: 0.0797 - val_accuracy: 0.9691
Epoch 153/200
10/10 [=====] - 0s 38ms/step - loss: 0.0554 - accuracy: 0.9845 - val_loss: 0.0734 - val_accuracy: 0.9725
Epoch 154/200
10/10 [=====] - 0s 36ms/step - loss: 0.0583 - accuracy: 0.9811 - val_loss: 0.0708 - val_accuracy: 0.9725
Epoch 155/200
10/10 [=====] - 0s 36ms/step - loss: 0.0483 - accuracy: 0.9871 - val_loss: 0.0644 - val_accuracy: 0.9759
Epoch 156/200
10/10 [=====] - 0s 36ms/step - loss: 0.0538 - accuracy: 0.9793 - val_loss: 0.0647 - val_accuracy: 0.9759
Epoch 157/200
10/10 [=====] - 0s 39ms/step - loss: 0.0526 - accuracy: 0.9793 - val_loss: 0.0629 - val_accuracy: 0.9828
Epoch 158/200
10/10 [=====] - 0s 35ms/step - loss: 0.0378 - accuracy: 0.9888 - val_loss: 0.0601 - val_accuracy: 0.9828
Epoch 159/200
10/10 [=====] - 0s 36ms/step - loss: 0.0533 - accuracy: 0.9828 - val_loss: 0.0630 - val_accuracy: 0.9725
Epoch 160/200
10/10 [=====] - 0s 36ms/step - loss: 0.0488 - accuracy: 0.9785 - val_loss: 0.0588 - val_accuracy: 0.9828
Epoch 161/200
10/10 [=====] - 0s 38ms/step - loss: 0.0412 - accuracy: 0.9845 - val_loss: 0.0731 - val_accuracy: 0.9759
Epoch 162/200
10/10 [=====] - 0s 37ms/step - loss: 0.0438 - accuracy: 0.9845 - val_loss: 0.0801 - val_accuracy: 0.9759
Epoch 163/200
10/10 [=====] - 0s 39ms/step - loss: 0.0363 - accuracy: 0.9871 - val_loss: 0.0774 - val_accuracy: 0.9725
Epoch 164/200
10/10 [=====] - 0s 36ms/step - loss: 0.0371 - accuracy: 0.9871 - val_loss: 0.0693 - val_accuracy: 0.9725
Epoch 165/200
10/10 [=====] - 0s 36ms/step - loss: 0.0436 - accuracy: 0.9811 - val_loss: 0.0738 - val_accuracy: 0.9725
Epoch 166/200
10/10 [=====] - 0s 36ms/step - loss: 0.0358 - accuracy: 0.9879 - val_loss: 0.0857 - val_accuracy: 0.9725
Epoch 167/200
10/10 [=====] - 0s 36ms/step - loss: 0.0484 - accuracy: 0.9811 - val_loss: 0.0839 - val_accuracy: 0.9759
Epoch 168/200
10/10 [=====] - 0s 39ms/step - loss: 0.0477 - accuracy: 0.9862 - val_loss: 0.0795 - val_accuracy: 0.9759
Epoch 169/200
10/10 [=====] - 0s 36ms/step - loss: 0.0489 - accuracy: 0.9854 - val_loss: 0.0935 - val_accuracy: 0.9656
Epoch 170/200
10/10 [=====] - 0s 40ms/step - loss: 0.0563 - accuracy: 0.9811 - val_loss: 0.0731 - val_accuracy: 0.9725
Epoch 171/200
10/10 [=====] - 0s 36ms/step - loss: 0.0445 - accuracy: 0.9836 - val_loss: 0.0642 - val_accuracy: 0.9759
Epoch 172/200
10/10 [=====] - 0s 36ms/step - loss: 0.0523 - accuracy: 0.9811 - val_loss: 0.0714 - val_accuracy: 0.9725
Epoch 173/200
10/10 [=====] - 0s 36ms/step - loss: 0.0664 - accuracy: 0.9811 - val_loss: 0.0920 - val_accuracy: 0.9691
Epoch 174/200
10/10 [=====] - 0s 36ms/step - loss: 0.0499 - accuracy: 0.9836 - val_loss: 0.0876 - val_accuracy: 0.9691
Epoch 175/200
10/10 [=====] - 0s 36ms/step - loss: 0.0371 - accuracy: 0.9819 - val_loss: 0.0760 - val_accuracy: 0.9725
Epoch 176/200
10/10 [=====] - 0s 40ms/step - loss: 0.0395 - accuracy: 0.9854 - val_loss: 0.0863 - val_accuracy: 0.9725
Epoch 177/200
10/10 [=====] - 0s 36ms/step - loss: 0.0458 - accuracy: 0.9836 - val_loss: 0.0874 - val_accuracy: 0.9691
Epoch 178/200
10/10 [=====] - 0s 35ms/step - loss: 0.0333 - accuracy: 0.9879 - val_loss: 0.0901 -


```
val_accuracy: 0.9691
Epoch 179/200
10/10 [=====] - 0s 37ms/step - loss: 0.0386 - accuracy: 0.9888 - val_loss: 0.0773 -
val_accuracy: 0.9691
Epoch 180/200
10/10 [=====] - 0s 36ms/step - loss: 0.0340 - accuracy: 0.9879 - val_loss: 0.0732 -
val_accuracy: 0.9759
Epoch 181/200
10/10 [=====] - 0s 36ms/step - loss: 0.0347 - accuracy: 0.9888 - val_loss: 0.0677 -
val_accuracy: 0.9828
Epoch 182/200
10/10 [=====] - 0s 36ms/step - loss: 0.0358 - accuracy: 0.9871 - val_loss: 0.0669 -
val_accuracy: 0.9794
Epoch 183/200
10/10 [=====] - 0s 36ms/step - loss: 0.0376 - accuracy: 0.9845 - val_loss: 0.0680 -
val_accuracy: 0.9828
Epoch 184/200
10/10 [=====] - 0s 36ms/step - loss: 0.0633 - accuracy: 0.9802 - val_loss: 0.0687 -
val_accuracy: 0.9725
Epoch 185/200
10/10 [=====] - 0s 40ms/step - loss: 0.0390 - accuracy: 0.9854 - val_loss: 0.0801 -
val_accuracy: 0.9725
Epoch 186/200
10/10 [=====] - 0s 36ms/step - loss: 0.0371 - accuracy: 0.9871 - val_loss: 0.0804 -
val_accuracy: 0.9759
Epoch 187/200
10/10 [=====] - 0s 39ms/step - loss: 0.0342 - accuracy: 0.9897 - val_loss: 0.0792 -
val_accuracy: 0.9794
Epoch 188/200
10/10 [=====] - 0s 36ms/step - loss: 0.0384 - accuracy: 0.9905 - val_loss: 0.0739 -
val_accuracy: 0.9863
Epoch 189/200
10/10 [=====] - 0s 36ms/step - loss: 0.0440 - accuracy: 0.9836 - val_loss: 0.0728 -
val_accuracy: 0.9828
Epoch 190/200
10/10 [=====] - 0s 37ms/step - loss: 0.0425 - accuracy: 0.9871 - val_loss: 0.0683 -
val_accuracy: 0.9828
Epoch 191/200
10/10 [=====] - 0s 40ms/step - loss: 0.0255 - accuracy: 0.9931 - val_loss: 0.0740 -
val_accuracy: 0.9863
Epoch 192/200
10/10 [=====] - 0s 36ms/step - loss: 0.0305 - accuracy: 0.9854 - val_loss: 0.0763 -
val_accuracy: 0.9794
Epoch 193/200
10/10 [=====] - 0s 36ms/step - loss: 0.0385 - accuracy: 0.9897 - val_loss: 0.0719 -
val_accuracy: 0.9863
Epoch 194/200
10/10 [=====] - 0s 36ms/step - loss: 0.0425 - accuracy: 0.9862 - val_loss: 0.0714 -
val_accuracy: 0.9828
Epoch 195/200
10/10 [=====] - 0s 36ms/step - loss: 0.0517 - accuracy: 0.9836 - val_loss: 0.0925 -
val_accuracy: 0.9691
Epoch 196/200
10/10 [=====] - 0s 35ms/step - loss: 0.0364 - accuracy: 0.9862 - val_loss: 0.0909 -
val_accuracy: 0.9725
Epoch 197/200
10/10 [=====] - 0s 36ms/step - loss: 0.0325 - accuracy: 0.9879 - val_loss: 0.0829 -
val_accuracy: 0.9828
Epoch 198/200
10/10 [=====] - 0s 35ms/step - loss: 0.0277 - accuracy: 0.9897 - val_loss: 0.0762 -
val_accuracy: 0.9828
Epoch 199/200
10/10 [=====] - 0s 36ms/step - loss: 0.0250 - accuracy: 0.9914 - val_loss: 0.0731 -
val_accuracy: 0.9863
Epoch 200/200
10/10 [=====] - 0s 36ms/step - loss: 0.0350 - accuracy: 0.9871 - val_loss: 0.0734 -
val_accuracy: 0.9863
```


Plotting Loss Values and Accuracy against Number of Epochs for Train Set and Test Set

In [16]:

```
figure = plt.figure(figsize=(10, 10))
plt.plot(hist.history['accuracy'], label='Train_accuracy')
plt.plot(hist.history['val_accuracy'], label='Test_accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc="upper left")
plt.show()

figure2 = plt.figure(figsize=(10, 10))
plt.plot(hist.history['loss'], label='Train_loss')
plt.plot(hist.history['val_loss'], label='Test_loss')
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc="upper left")
plt.show()
```





Evaluating Model on Test Set

```
In [17]: pred = model.evaluate(X_test, Y_test)
print(f'Test Set Accuracy: {pred[1]}')
print(f'Test Set Loss: {pred[0]}')
```

10/10 [=====] - 0s 7ms/step - loss: 0.0734 - accuracy: 0.9863
Test Set Accuracy: 0.9862542748451233
Test Set Loss: 0.07337489724159241

Classification Report


```
In [18]: ypred = model.predict(X_test)
ypred = np.argmax(ypred, axis=1)
Y_test_pred = np.argmax(Y_test, axis=1)
print(classification_report(Y_test_pred, ypred))
```

	precision	recall	f1-score	support
0	0.99	0.98	0.99	169
1	0.98	0.99	0.98	122
accuracy			0.99	291
macro avg	0.98	0.99	0.99	291
weighted avg	0.99	0.99	0.99	291

Confusion Matrix

In [19]:

```
matrix = confusion_matrix(Y_test_pred, ypred)
df_cm = pd.DataFrame(matrix, index=[0, 1], columns=[0, 1])
figure = plt.figure(figsize=(5, 5))
sns.heatmap(df_cm, annot=True, fmt='d')
```



Saving the Model

In [20]:

```
model.save('Driver_Drowsiness_Detection.h5')
```

Testing Model on Images

Testing the classification performnace on random images of Open and Closed Eyes

In [21]:

```
labels = ['Closed', 'Open']
img_closed1 = cv2.imread('driver_drowsiness_detection/closed_eye.jpg')
img_closed2 = cv2.imread('driver_drowsiness_detection/closed_eye2.jpg')
img_open1 = cv2.imread('driver_drowsiness_detection/open_eye.jpg')
img_open2 = cv2.imread('driver_drowsiness_detection/open_eye2.jpg')

img_closed1 = cv2.resize(img_closed1, (32, 32))
img_closed2 = cv2.resize(img_closed2, (32, 32))
img_open1 = cv2.resize(img_open1, (32, 32))
img_open2 = cv2.resize(img_open2, (32, 32))

img_closed1 = np.array(img_closed1)
img_closed2 = np.array(img_closed2)
img_open1 = np.array(img_open1)
img_open2 = np.array(img_open2)

img_closed1 = np.expand_dims(img_closed1, axis=0)
img_closed2 = np.expand_dims(img_closed2, axis=0)
img_open1 = np.expand_dims(img_open1, axis=0)
img_open2 = np.expand_dims(img_open2, axis=0)
```

In [22]:

```
ypred_closed1 = model.predict(img_closed1)
ypred_closed2 = model.predict(img_closed2)
ypred_open1 = model.predict(img_open1)
ypred_open2 = model.predict(img_open2)
```

In [23]:

```
figure = plt.figure(figsize=(2, 2))
img_closed1 = np.squeeze(img_closed1, axis=0)
plt.imshow(img_closed1)
plt.axis('off')
plt.title(f'Prediction by the model: {labels[np.argmax(ypred_closed1[0], axis=0)]}')
plt.show()
```


Prediction by the model: Closed



In [24]:

```
figure = plt.figure(figsize=(2, 2))
img_closed2 = np.squeeze(img_closed2, axis=0)
plt.imshow(img_closed2)
plt.axis('off')
plt.title(f'Prediction by the model: {labels[np.argmax(ypred_closed2[0], axis=0)]}')
plt.show()
```

Prediction by the model: Closed



In [25]:

```
figure = plt.figure(figsize=(2, 2))
img_open1 = np.squeeze(img_open1, axis=0)
plt.imshow(img_open1)
plt.axis('off')
plt.title(f'Prediction by the model: {labels[np.argmax(ypred_open1[0], axis=0)]}')
plt.show()
```

Prediction by the model: Open



In [26]:

```
figure = plt.figure(figsize=(2, 2))
img_open2 = np.squeeze(img_open2, axis=0)
plt.imshow(img_open2)
plt.axis('off')
plt.title(f'Prediction by the model: {labels[np.argmax(ypred_open2[0], axis=0)]}')
plt.show()
```

Prediction by the model: Open



Creating Pipeline for making predictions on full face images


In [27]:

```
def full_face_detection_pipeline(input_image_path):
    face_cascade = cv2.CascadeClassifier('driver_drowsiness_detection/haarcascade_frontalface_default.xml')
    eye_cascade = cv2.CascadeClassifier('driver_drowsiness_detection/haarcascade_eye.xml')
    detector = dlib.get_frontal_face_detector()
    predictor = dlib.shape_predictor('driver_drowsiness_detection/shape_predictor_68_face_landmarks.dat')
    fa = FaceAligner(predictor, desiredFaceWidth=256)
    test_image = cv2.imread(input_image_path)
    test_image = imutils.resize(test_image, width=800)
    test_image_gray = cv2.cvtColor(test_image, cv2.COLOR_BGR2GRAY)
    rects = detector(test_image_gray, 2)
    for rect in rects:
        (x, y, w, h) = rect_to_bb(rect)
        faceOrig = imutils.resize(test_image[y:y+h, x:x+w], width=256)
        faceAligned = fa.align(test_image, test_image_gray, rect)
        faceAligned_gray = cv2.cvtColor(faceAligned, cv2.COLOR_BGR2GRAY)
        plt.imshow(faceAligned_gray)
        plt.axis('off')
        plt.title('Aligned Face')
        plt.show()
        eyes = eye_cascade.detectMultiScale(faceAligned_gray, 1.1, 4)
        predictions = []
        for (ex, ey, ew, eh) in eyes:
            eye = faceAligned[ey:ey+eh, ex:ex+ew]
            # cv2.rectangle(test_image, (x+ex, y+ey), (x+ex+ew, y+ey+eh), (0, 0, 255), 8)
            eye = cv2.resize(eye, (32, 32))
            eye = np.array(eye)
            eye = np.expand_dims(eye, axis=0)
            ypred = model.predict(eye)
            ypred = np.argmax(ypred[0], axis=0)
            predictions.append(ypred)
        if all(i==0 for i in predictions):
            cv2.rectangle(test_image, (x, y), (x+w, y+h), (0, 0, 255), 8)
            cv2.putText(test_image, 'Sleeping', (x, y-10), cv2.FONT_HERSHEY_SIMPLEX, 1.5, (0, 0, 255), 3)
        else:
            cv2.rectangle(test_image, (x, y), (x+w, y+h), (0, 255, 0), 8)
            cv2.putText(test_image, 'Not Sleeping', (x, y-10), cv2.FONT_HERSHEY_SIMPLEX, 1.5, (0, 255, 0), 3)
    output_path = 'driver_drowsiness_detection/test_image_prediction.jpg'
    cv2.imwrite(output_path, test_image)
    return output_path
```

In [28]:

```
figure = plt.figure(figsize=(5, 5))
predicted_image = cv2.imread(full_face_detection_pipeline('driver_drowsiness_detection/active_person.jpg'))
predicted_image = cv2.cvtColor(predicted_image, cv2.COLOR_BGR2RGB)
plt.imshow(predicted_image)
plt.axis('off')
plt.show()
```

Aligned Face



In [29]:

```
figure = plt.figure(figsize=(5, 5))
predicted_image = cv2.imread(full_face_detection_pipeline('driver_drowsiness_detection/drowsy_person.jpg'))
predicted_image = cv2.cvtColor(predicted_image, cv2.COLOR_BGR2RGB)
plt.imshow(predicted_image)
plt.axis('off')
plt.show()
```


Aligned Face



In [30]:

```
figure = plt.figure(figsize=(5, 5))
predicted_image = cv2.imread(full_face_detection_pipeline('driver_drowsiness_detection/sleepy-driver.jpg'))
predicted_image = cv2.cvtColor(predicted_image, cv2.COLOR_BGR2RGB)
plt.imshow(predicted_image)
plt.axis('off')
plt.show()
```

Aligned Face



In []: